

Installing Content-Aware Storage (CAS)

Last Updated: 2026-04-01

Initiate the installation of CAS from the user interface of IBM Fusion. It installs Kafka, CAS operator, and CNPG.

Before you begin [↗](#)

- Proceed with the installation only after all the prerequisites are met. For more information about prerequisites, see [Planning and prerequisites](#).
- Ensure that you have storage class and is marked as default. For example, `ibm-spectrum-fusion`. For more information about how to mark a storage class as default, see [Prerequisites](#) topic.
- Ensure all Scale projects and pods are running.
- For IBM Fusion services and platform support matrix, see [Support matrix for IBM Fusion versions](#).

About this task [↗](#)

Use this procedure to install the CAS service from the Services page of the IBM Fusion user interface. You can also install the CAS service using CLI. For more information about CLI usage, see [Configuring Content-Aware Storage \(CAS\) through CLI](#).

Follow the procedure if:

- Your data resides on an external IBM Storage Scale filesystem.
- You prefer to use an external IBM Storage Scale filesystem as cache filesystem for your S3 data source.

Otherwise, follow the procedure in [Installing Content-Aware Storage \(CAS\) without a remotely mounted IBM Storage Scale filesystem](#).

Procedure [↗](#)

1. In the IBM Fusion user interface, go to Services page.
2. In the **Available** section, click the **CAS** tile.
3. In the **CAS** page, go through the features and capabilities of the service and click **Install**.

If the pre-check raises a warning, then resolve it at OpenShift level.

The installation starts and a notification appears on the Services page. You can see the progress of the installation in the Services > **Installed** section. After the installation is completed successfully, you can see the status as **Healthy**. A success notification gets displayed. It disappears automatically after some time.

From the ellipsis menu, you can download logs and view documentation.

4. In the Install service window, select storage class and click **Install**.
5. Validate the installation.
 - Verify the CAS installation from the IBM Fusion user interface:

After you enable the CAS service, you can view the service version and health status. From the ellipsis menu, you can download logs and view documentation. After you successfully collect the logs, a success notification is displayed. The notification disappears automatically after some time.
 - Verify the CAS installation from the OpenShift® Container Platform web console:
 - a. In the OpenShift console, go to **Home** > **Projects**.
 - b. In the Projects page, search and check for the availability of a namespace with string `ibm-cas`.
 - c. Go to **Operators** > **Installed Operators** and check for the availability of CAS.
 - d. Open **Content Aware Storage** and check whether a CAS service definition is available.
 - e. After you enable the CAS service, you can view the service version and health status. From the ellipsis menu, you can download logs and view documentation. After you successfully collect the logs, a success notification is displayed. The notification disappears automatically after some time.

What to do next

- To verify the percentage completion of prerequisites post installation, do the following steps:
 1. Open IBM Fusion user interface.
 2. Go to **Content-Aware Storage** > **Overview**.
 3. In the **Jump right in** section, click **Getting started with CAS** tile.
 4. A side panel opens with the status of the following prerequisites:
 - Integrate NVIDIA NIM
 - Set up of Kafka Broker
 - Connect caching file system
- After you successfully install CAS, set up a Scale user to enable watch creation, configure domains and data sources. Map the domain to the data source. For more information about these procedures, see [Configuring Content-Aware Storage \(CAS\)](#).
- **Planning and prerequisites**

Go through the Content-Aware Storage (CAS) installation prerequisites before you install the service.
- **[Installing Content-Aware Storage \(CAS\) without a remotely mounted IBM Storage Scale filesystem](#)**
- **[Installing NVIDIA NeMo Retriever Extraction \(nv-ingest\)](#)**

The Content-Aware Storage (CAS) Document Processor Engine can either use IBM Docling Multimodal or NVIDIA NeMo Retriever Extraction for document extraction.
- **[Configuring NVIDIA Inference Microservices \(NIMs\)](#)**

After you install CAS, the `cas-config` configmap must be created. It provides the `nv-ingest` endpoint information.

- [Configuring a Multi-Instance GPU \(MIG\) with CAS](#)

Multi-Instance GPU (MIG) partitions a single NVIDIA GPU into multiple isolated instances. Each instance operates as an independent GPU, enabling resource sharing across multiple deployments.

- [Configuring IBM Storage Scale user to enable watch creation](#)

Configure an IBM Storage Scale user with all the needed privileges to enable watch creation.

Parent topic:

→ [Content-Aware Storage \(CAS\)](#)